Experiments in Integrated Data Center Power and Cooling Analysis

Shaun Brandt, Hema Kumar, Michael Novak, Mandar Patil, Yu Yang

Portland State University

About the Project

- The need for large scale computing continues to rise, and companies are turning to data centers to meet those needs
- Large clusters of servers are expensive to run, and expensive to cool
- Cooling devices (especially chillers) consume lots of power and are subject to physical wear
- Data center operators need access to as much information as possible!

About the Data

- Data for this project was collected from a number of sensors at the PNNL Energy Smart Data Center
- 8 server racks (1 network, 2 air cooled, 5 liquid cooled)
- Two application loads (high density / low density)
- Three different application test runs

Project Design

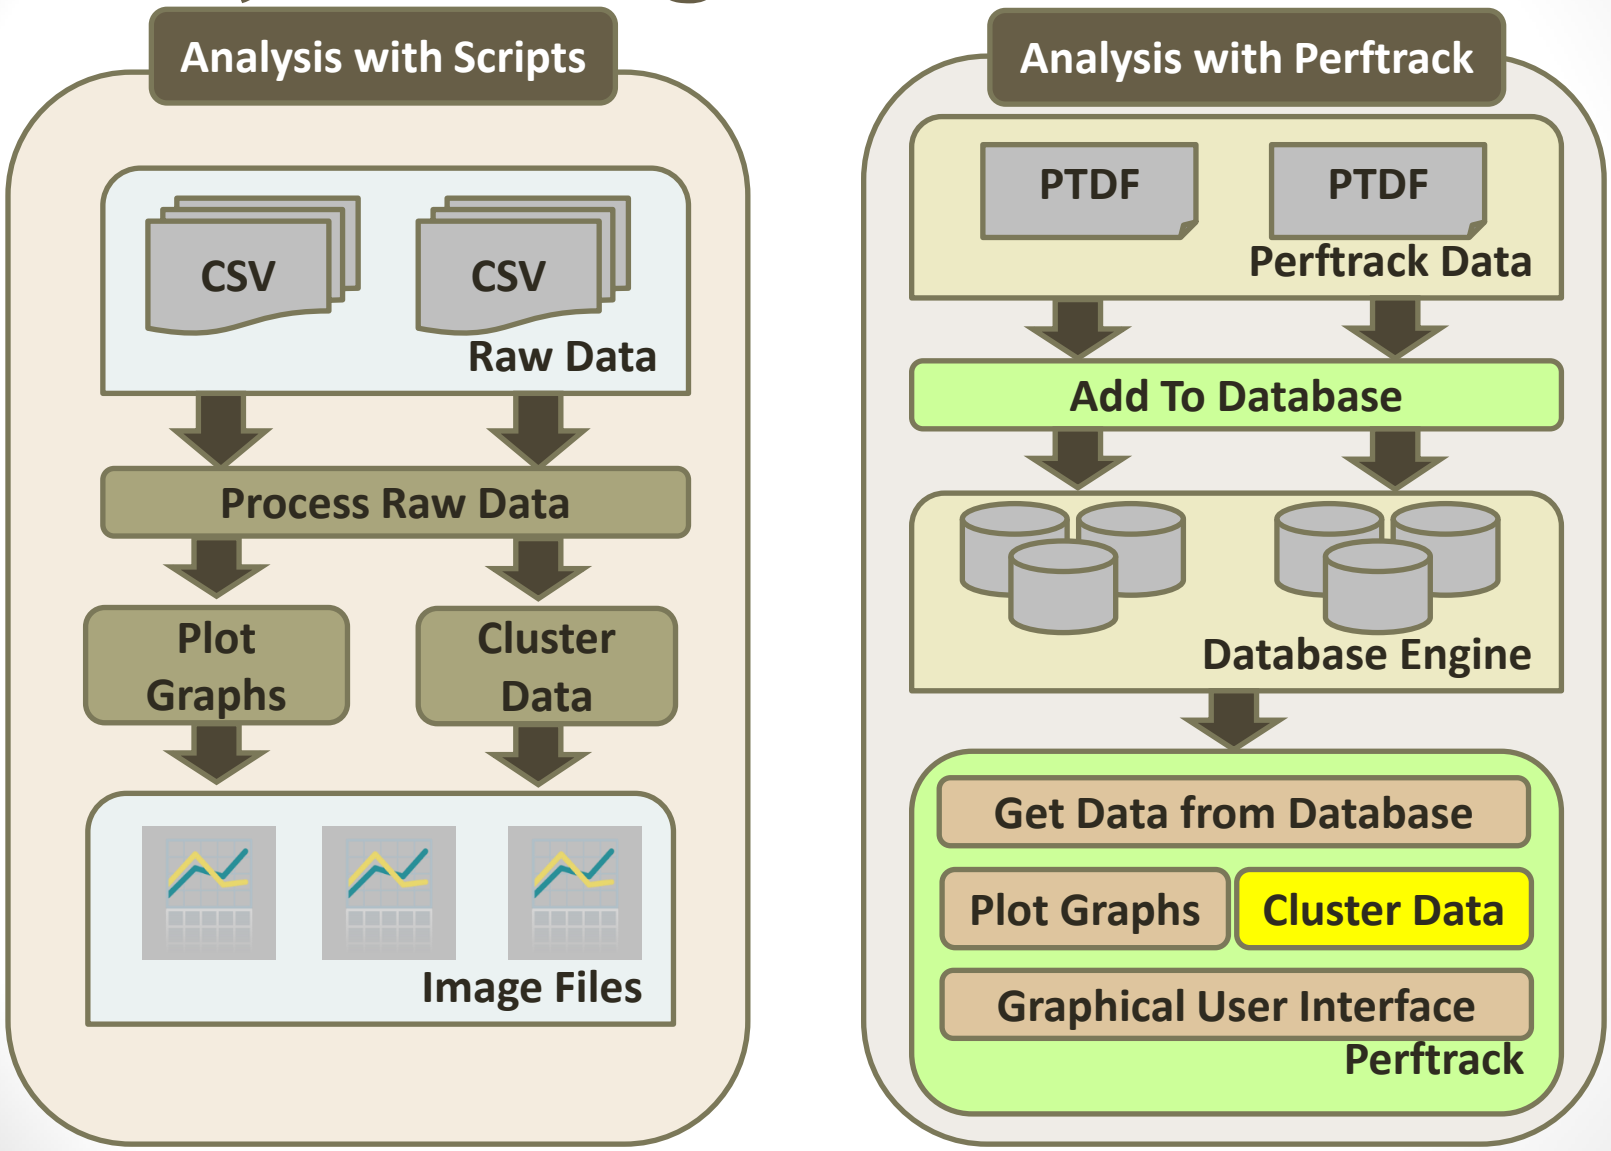

Approach 1: Analysis with Scripts

- A set of custom tools was written to work with sensor data provided in CSV format
- Mostly Haskell with calls to Gnuplot (earlier prototypes also used C, Perl and shell)
- Generates a complete report including tables and graphs
- Calculates and graphs metrics such as temperature deltas, COP (coefficient of performance) and power usage

Analysis with Scripts (cont.)

- Cluster analysis based on output of metric calculations
- Our work uses ELKI, an open source data mining tool with many customizable options (such as choice of clustering algorithms)
- We tried clustering on various metrics, including chiller water temperature deltas, coefficient of performance (COP) and cooling load
- We developed a method to overlay cluster information over previously generated graphs

Time-correlated Cluster Charts

- On the left – low density chiller temperature delta chart
- On the right – the associated time-correlated chart

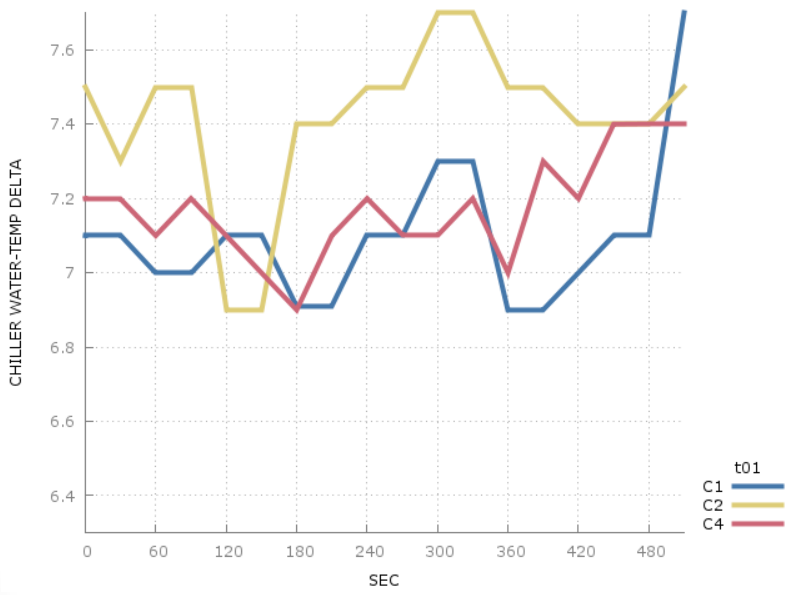

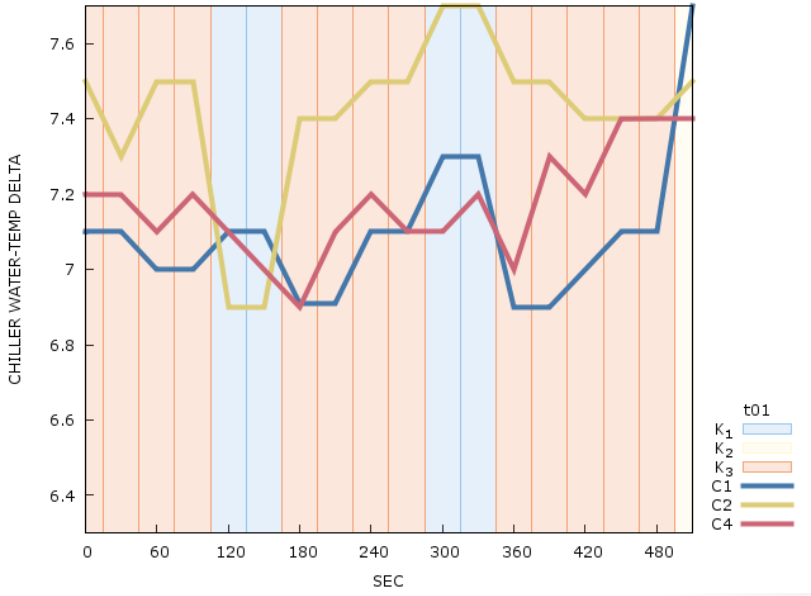

Chiller Power (KW) Chart, HD Data

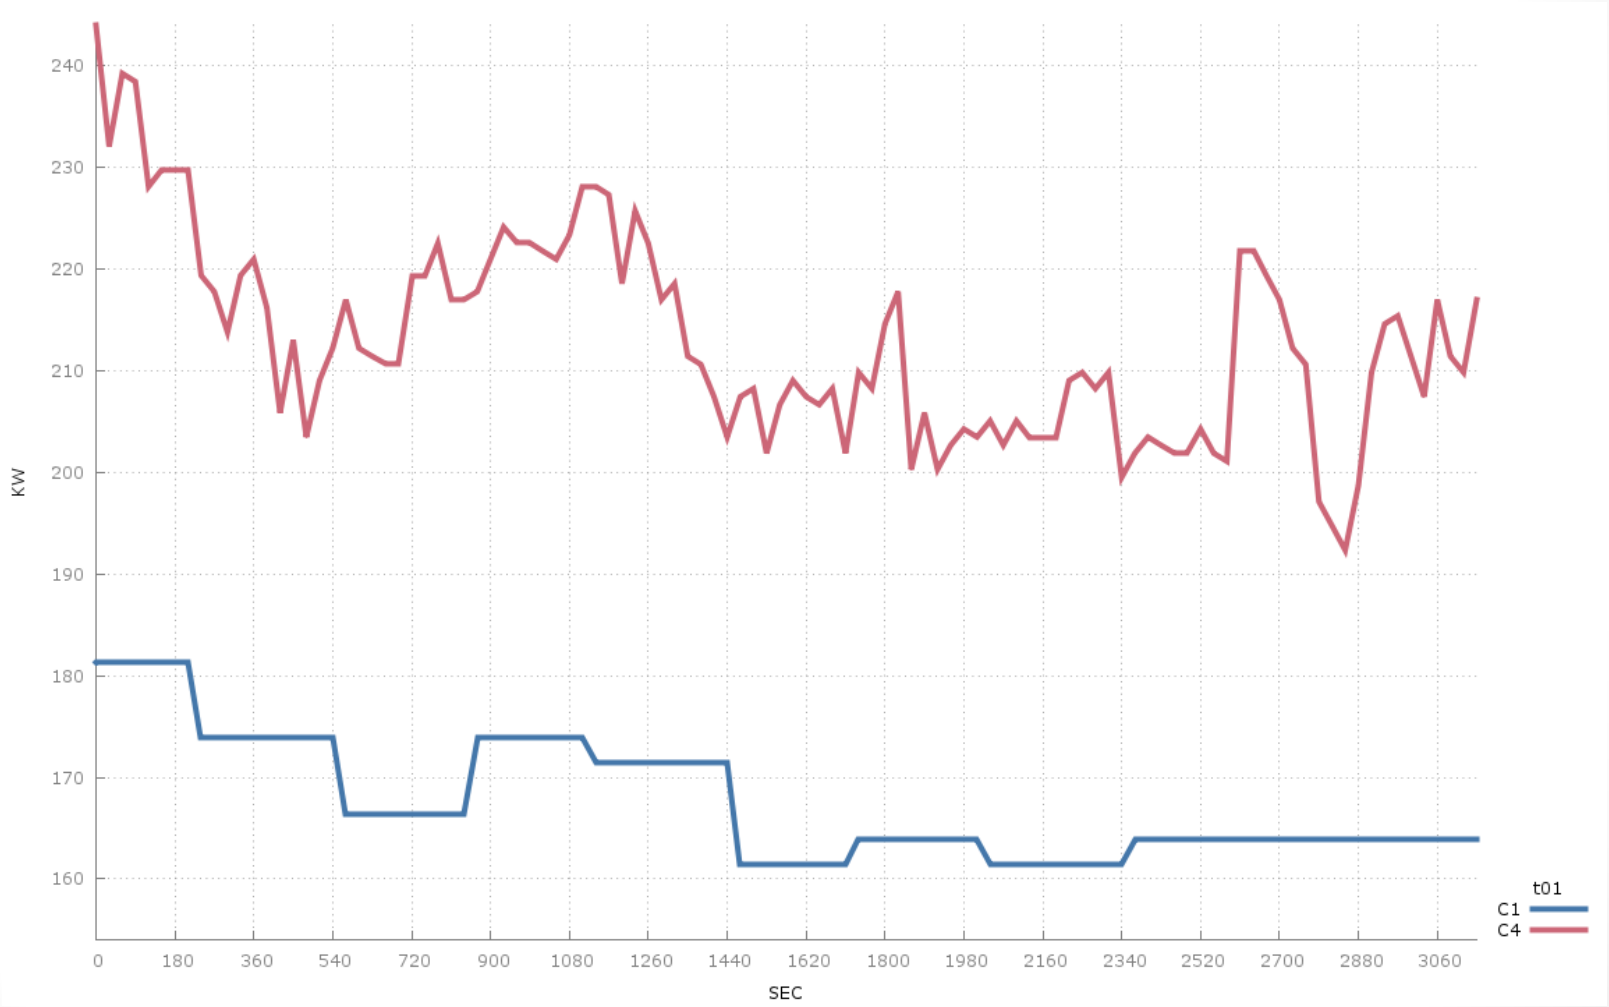

Associated Time-correlated Cluster Chart

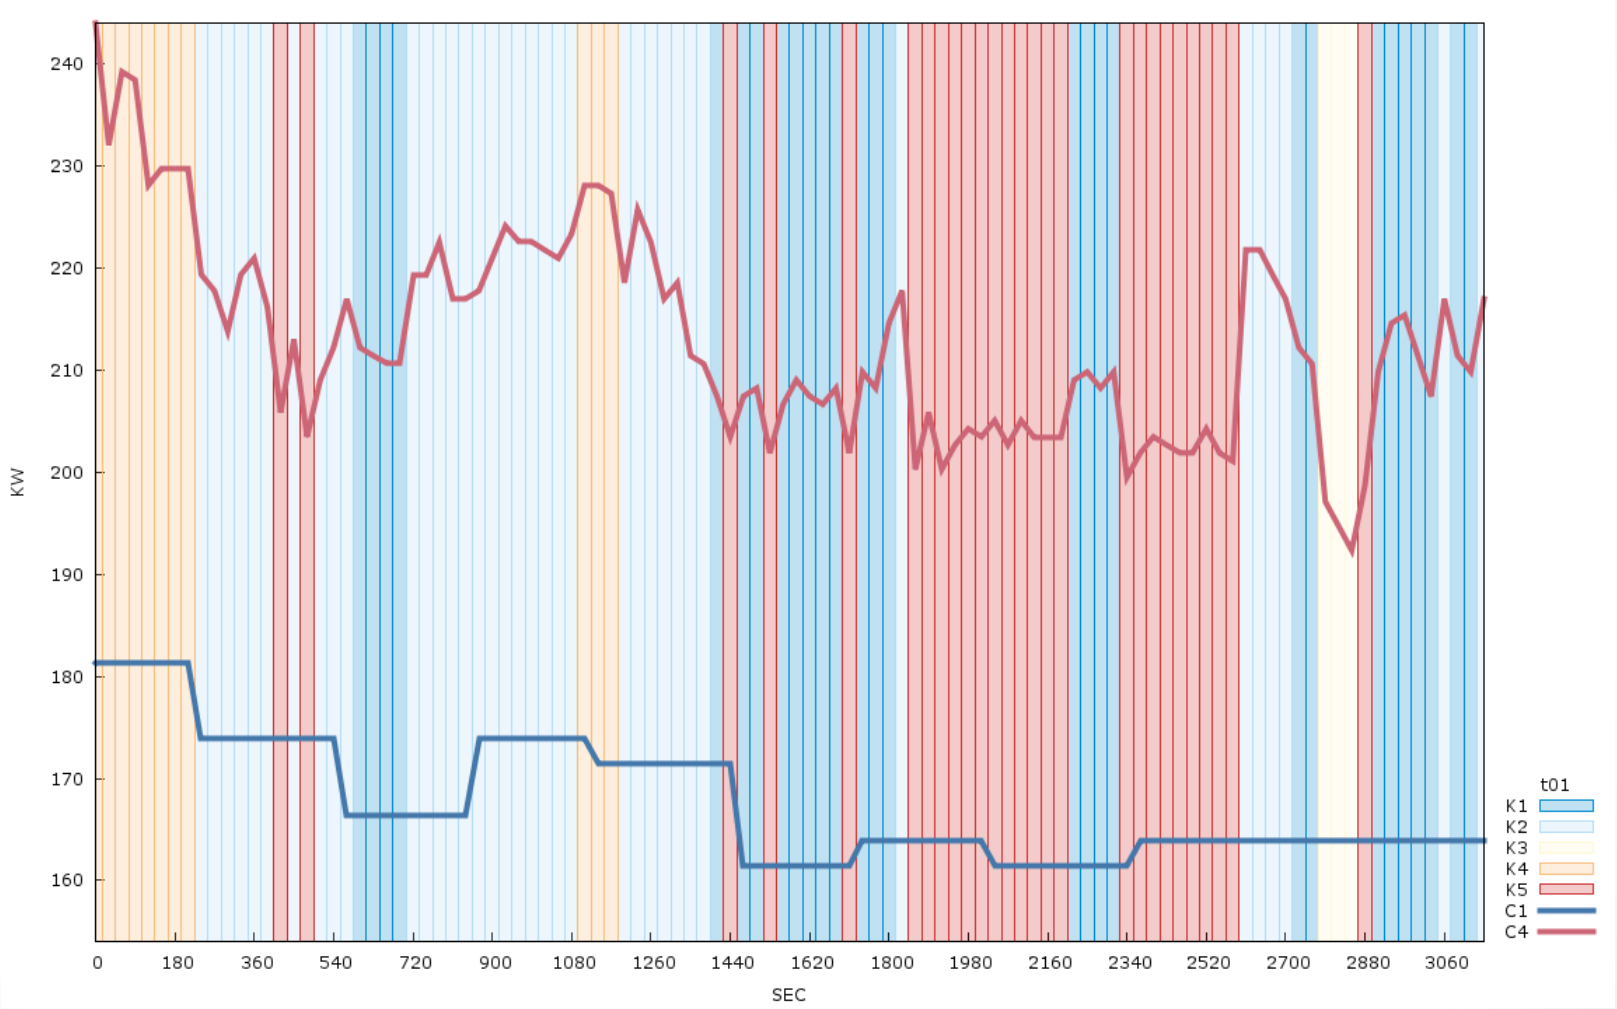

Perftrack: Data Analysis Tool

- Data store and interface for managing performance data
- Compares the data collected in different locations and formats in single performance analysis session
- Includes interfaces to data store and scripts for automatically collecting data describing each experiment
- Includes a graphical user interface
- Written in C++, QT and Python

Approach 2: Perftrack Modifications

- ELKI: Open source, Java-based clustering package
- New menu item ‘Cluster Data’ is added to Perftrack
- Perftrack is modified to generate the ELKI-compatible raw data from user filtered experiment data using GUI
- K-Means clustering algorithm is executed on ELKI-compatible raw data using ELKI package
- Perftrack displays the clusters

Future Work

- Need more (and more accurate) data!
- Perftrack:
  - Full integration of script work
  - On-the-fly data generation for ELKI
  - Tightly-coupled integration of ELKI
  - Graphical user interface enhancement
- Modification of scripts to support direct connection to FRED in place of CSV file access
- Experiment with motif-based cluster analysis (this would require more data)